

Chapter 4: Crease-Aware Subgradient Selection

(Experiment 1)

A Technical Report from the Puno Calculus Research Program

Michael Grafiel Sayson Puno | July 2026

Question

Does the choice of subgradient at ReLU's non-differentiable point ($z = 0$) affect generalization?

This is the most direct test of the Puno Calculus claim: if what happens at the crease matters, then different choices at the crease should lead to different outcomes.

Setup

- **Problem:** Binary classification on concentric ring dataset
- **Architecture:** MLP [2, 64, 64, 1]
- **Training:** Adam, 300 epochs, batch size 128
- **Four subgradient strategies at $\text{ReLU}(z = 0)$:**
- **Standard:** $\partial \text{ReLU} / \partial z = 0$ (PyTorch/TensorFlow default)
- **Random:** $\partial \text{ReLU} / \partial z = \text{Bernoulli}(0.5)$ at each crease crossing
- **Always-on:** $\partial \text{ReLU} / \partial z = 1$ (treats $z = 0$ as in the active region)
- **Oppose:** $\partial \text{ReLU} / \partial z = -1$ (flip gradient at crease)

Results

Strategy	Test Accuracy	Final Loss	Crease Density (mean/step)
Standard (0)	97.52%	0.1058	152.3
Random	97.55%	0.1053	157.2
Always-on (1)	97.67%	0.1055	**118.4**
Oppose	97.48%	**0.1026**	155.8

Key Findings

1. **All strategies converge to similar accuracy.** On simple problems, subgradient choice at the crease does not affect final performance. The network finds its way regardless.

2. **BUT: the crease strategy changes how the network learns.** The "always-on" strategy produced systematically fewer near-crease units during training (~118/step vs ~152-157 for others). This means the subgradient choice at $z = 0$ affects whether units stay near their switching threshold or move decisively into on/off states.

3. **The "oppose" strategy achieves slightly lower final loss** (0.1026 vs 0.1058). Flipping the gradient at crease boundaries appears to help the optimization settle into a slightly better minimum.

Interpretation

The Puno Calculus predicts that crease behavior is consequential. Experiment 1 confirms this prediction but reveals an important nuance: **on simple problems, the effect is small.** The network has enough capacity and the problem is easy enough that any strategy works.

This is consistent with the theoretical result from Bertoin et al. (NeurIPS 2021): in infinite precision, the bifurcation zone where subgradient choice matters is measure-zero. At finite precision (we used float64), the effect is real but small for easy problems.

The key open question: does crease-aware subgradient selection matter more on *hard* problems — problems where the network is pushed to its capacity limits? Our experiments cannot answer this, but the literature suggests yes.

Literature Connection

Boursier et al. (2022, NeurIPS) proved a result stronger than what we previously reported. The paper shows that **only the subgradient $\sigma'(x) = 1_{\{x > 0\}}$ ($\text{ReLU}'(0) = 1$) guarantees the existence of a globally defined gradient flow solution.** Any other constant choice — including $\text{ReLU}'(0) = 0$, the standard in PyTorch and TensorFlow — may cause the gradient flow ODE to fail to have a globally defined solution. Trajectories can blow up or cease to exist.

This is a mathematical existence result about the continuous-time dynamics: the subgradient choice determines whether the optimization is even *well-posed*. The paper further shows that gradient flow with $\text{ReLU}'(0) = 1$ converges to the minimum L2-norm interpolator on orthogonal data with vanishing initialization. The variation norm — the total variation of the second derivative of the learned function — is the implicit bias.

The practical implication (consistent with Bertoin et al., NeurIPS 2021): $\text{ReLU}'(0) = 0$ is standard in frameworks for good empirical reasons (it works better with SGD at finite precision), but the continuous-time theory demands $\text{ReLU}'(0) = 1$. This tension between theory and practice is exactly where the Puno Calculus operates: the crease is consequential, and understanding how it works requires looking at both regimes.

The Boursier paper is the single most important theoretical result for the Puno Calculus: different subgradient choices don't just change the trajectory — they can determine whether a solution exists at all.
